

Toward building a benchmark for LLM business-oriented cross-cultural understanding

Anonymous ACL submission

Abstract

Large Language Models (LLMs) have improved rapidly, but alignment research still focuses mainly on universal values, often overlooking cultural diversity. This work introduces the LLM Cultural Map Benchmark, which evaluates cross-cultural awareness across Erin Meyer’s eight behavioural dimensions.

The benchmark addresses two common validity problems: memorisation of public cultural sources and answer-position bias. It includes B1, with 320 comparative Yes/No questions per variant (40 per dimension), B2, with 160 applied workplace scenarios per variant (20 per dimension), and B3, with 80 behaviour-identification items. B1 and B2 each ship in two wording variants (A and B) as a built-in format-robustness control.

We also map Culture-Conflict Decision Benchmark (CCD-Bench) and EtiCor++ to Meyer’s framework using rule-based, LLM-assisted, and human validation. Around 79% of CCD-Bench items do not align with Meyer’s dimensions. The benchmark covers ten Global Leadership and Organizational Behavior Effectiveness (GLOBE) clusters and includes 35% counter-stereotype items in B1 and B2 to discourage regional shortcut learning.

1 Introduction

Large Language Models (LLMs) are becoming more and more entangled in our daily lives, even though controlling their responses is becoming increasingly important. Focusing on cross-cultural professional contexts, we needed to ensure that the **WEIRD** (Western, Educated, Industrialized, Rich, and Democratic) **bias** (Peter and Devlin, 2025) is being controlled. WEIRD bias can happen either because of the training data, mostly web scraping of websites owned by entities from the WEIRD cultures (?), or because of cognitive homogenization (Sourati et al., 2026) and **RLHF** (Reinforcement Learning from Human Feedback) (Poddar et al.,

2024) is done by users of these LLMs, who are mostly represented by people from the WEIRD country-based cultures (Santurkar et al., 2023). The real-life harms that can happen from this kind of bias are automated content moderation, exemplified by the erroneous identification of non-WEIRD content (such as African-American Vernacular English [AAVE]) as "toxic", "harmful", or "offensive" (Sap et al., 2019), and poor performance for non-WEIRD users based on shallow stereotyping (Liu, 2024).

Our focus is specifically on **business-related communication**: the workplace is where cross-cultural misreadings carry immediate practical cost — misjudged feedback, stalled negotiations, misread silence — and it is the domain that Erin Meyer’s *The Culture Map* (Meyer, 2014) describes through **eight behavioural dimensions** (Communicating, Evaluating, Persuading, Leading, Deciding, Trusting, Disagreeing, and Scheduling) of professional interaction. The overall purpose of this work is to build a benchmark, the **LLM Cultural Map Benchmark**, that evaluates LLMs’ cross-cultural awareness in this business domain, grounded in Meyer’s framework, using the book’s abstract continuous scales representation of countries into quantifiable distance metrics to generate objective ground-truth data.

Before creating new datasets, we carefully explored existing resources in search of appropriately labelled data. Two candidates could have made our case: **CCD-Bench** (Rahman and Salam, 2026), which probes cultural conflict in LLM decision-making through open-ended dilemmas, and **EtiCor++** (Dwivedi et al., 2025), a large corpus of region-specific etiquette statements. However, our analysis showed that neither covers Meyer’s framework: 79% and 76% of their content, respectively, maps to no Meyer dimension, and the Persuading dimension is absent from both. This convinced us to create new evaluation datasets from scratch,

retaining CCD-Bench and EtiCor++ as external evidence of the gap — and as measured motivation for a purpose-built benchmark.

A second validity problem shaped the design: models can score highly (90+%) on cultural questions simply because the underlying sources — including some benchmarks’ own evaluation sets — are publicly available online and were seen during pre-training. To disentangle memorization (Morris et al., 2025) from understanding, our datasets consist of new questions built parametrically from facts in Meyer’s book (**B1** Yes/No factual questions and **B2** Multiple Choice Questions (MCQ) workplace scenarios), complemented by **B3** behaviour-identification items, so that the model sees each exact question and situation for the first time.

The remainder of this document is structured as follows. Sections 2–4 review related work and present our analysis of the two existing corpora, CCD-Bench and EtiCor++, whose limitations motivate the benchmark. Section 5 describes the construction of the new evaluation blocks (B1, B2 and B3). Section 6 reports the mapping validation results and the benchmark evaluation setup; Section 7 presents the evaluation results across four open models, and Sections 8–9 conclude and discuss limitations.

2 Related Work

Evaluating cross-cultural knowledge in Large Language Models has typically focused either on general opinion surveys, such as GlobalOpinionQA (Durmus et al., 2024), or on model cultural fine-tuning like CultureLLM (Li et al., 2024). While these approaches measure macro-level ideological distributions or try to inject regional data, they often fail to answer business-focused contexts. Furthermore, prior studies leveraging traditional cultural frameworks, like GLOBE (House et al., 2004), struggle with interaction-level nuances, normally leaving critical business-related dimensions unrepresented. In contrast, our project introduces a purpose-built benchmark based in Erin Meyer’s behavioral dimensions, using controlled workplace scenarios and counter-stereotype traps to audit cross-cultural competence without relying on static opinion retrieval.

3 CCD-Bench

3.1 Mapping and Validation

To assess whether CCD-Bench (Rahman and Salam, 2026) items could serve as evaluation questions, we filtered the 2,182 open-ended questions to 369 business/workplace related items by retaining only questions whose domain label matched one of seven work-context categories (Work, Scientific, Legal, Financial, Educational, Lifestyle, and Environmental), excluding purely personal or social questions. Then, we mapped its 369 unique question stems onto Meyer’s eight dimensions using a three-tier hybrid method: a rule-based lexical mapper, an LLM-assisted pass, and human validation as ground truth.

The rule-based mapper classifies each stem along with its answer options, defaults to Unmapped, and records the matched phrases for each label. It left 74.5% of the corpus Unmapped. The LLM pass (llama-3.3-70b-versatile, temperature 0, batched) classified the same content under a prompt with an explicit GLOBE-exclusion list. An ablation isolated the two changes introduced in this pass: making the answer options visible to the classifier recovers dimension assignment (+6 accuracy points), while hardening the prompt to explicitly exclude GLOBE constructs improves rejection of non-Meyer content (+3 accuracy points). A naive count of relabelled items would have reversed this attribution, suggesting the prompt change was more impactful; accuracy against ground truth shows the opposite.

Ground truth was built over three annotation rounds on 90 stratified items. Agreement went from $k = 0.31$ (blind, stems only) to $k = 0.56$ (blind, stems plus answer poles) with the same annotators and guideline, showing inter-annotator agreement is a function of the information available, not only of annotator skill. Against the final ground truth (68 consensus + 22 adjudicated items), the rules reach 57% and the final LLM mapping 78%. Note that because 79% of the corpus is Unmapped, a majority-class baseline would score approximately 79% overall; the 78% figure is therefore best interpreted as performance on the harder sub-task of correctly assigning the 21% of items that do have a Meyer dimension, where the LLM substantially outperforms the rules (57%). Evaluating only on consensus rows would have reported 88%; consensus rows are by construction the easy ones, so we report the adjudicated set.

3.2 Result: CCD-Bench Does Not Cover Meyer’s Dimensions

79% of filtered CCD-Bench (290 of 369) maps to no Meyer dimension, and human annotators independently agree (59 of 90 ground-truth items Unmapped). The corpus encodes GLOBE-style constructs — individualism–collectivism, long- and short-term orientation, uncertainty avoidance — for which Meyer’s framework has no corresponding dimensions. Persuading is absent under rules, LLM and humans alike, and six of the eight dimensions retain fewer than eight mapped items, making per-dimension scoring statistically meaningless. CCD-Bench therefore enters this work as external-corpus evidence rather than evaluation items: it demonstrates, with measured data, the gap that the new datasets presented in Section 5 are built to fill.

4 EtiCor++

4.1 Corpus, Mapping and Validation

EtiCor++ (Dwivedi et al., 2025) is a corpus of over 46,000 etiquette statements labelled by region. We restricted it to its business subset (11,264 statements) and applied the same rule-based lexicon used for CCD-Bench, unchanged, so the two outcomes are directly comparable.

The mapper leaves 76% Unmapped; among mapped items, Scheduling, Trusting, and Leading dominate, and Persuading is again absent — replicating the CCD result on a second, independently constructed corpus. Validation followed the same annotation protocol used for CCD-Bench (two independent annotators, pre-registered guidelines, blind annotation), yielding 91 consensus rows and $k = 0.80$; the rules reach 92% accuracy on the audited mapped items, monotonic across confidence tiers (high tier: 100%).

This is the decisive contrast with CCD-Bench: one lexicon, 92% on one corpus, 57% on the other. The lexical approach transfers to EtiCor’s declarative etiquette statements but fails on CCD’s introspective MCQs, a method-selection finding that explains why the LLM pass was adopted for CCD while rules were retained for EtiCor++.

4.2 Result: EtiCor++ As External Evidence

EtiCor++ is likewise used as external-corpus evidence rather than evaluation items, for three reasons: 76% of it has no Meyer dimension; its answers were set by the corpus authors with

a methodology we cannot audit; and it is validated at regional granularity, so assigning country-level answers would reintroduce the same-region-same-behaviour shortcut our counter-stereotype design exists to penalise (Japan and China, same macro-region, opposite poles of Deciding). An exploratory reformulation pipeline suggests only (~ 2.5 – 5)% of the corpus would survive as usable items — itself further evidence that etiquette corpora barely intersect Meyer’s interaction-level dimensions, and that purpose-built evaluation is required.

5 Dataset Construction: B1, B2 and B3

Having established that neither existing corpus covers Meyer’s framework, we built the evaluation datasets from scratch. Figure 1 shows the overall system pipeline, from the raw datasets through multi-model LLM inference to automated evaluation and reporting.

5.1 Motivation and Validity Problem

The benchmark consists of three evaluation blocks B1–B3, each testing a different level of cross-cultural reasoning.

- **B1:** presents comparative Yes/No questions between two countries on a specific Meyer dimension. Asks which one is more likely to exhibit a specific cultural behaviour.
- **B2 (Applied MCQ):** presents workplace misunderstanding scenarios and asks the model to identify the culturally grounded explanation among four options; testing whether dimensional knowledge transfers to applied reasoning.
- **B3 (Identification MCQ):** inverts the direction. The model reads an observable behaviour and must infer which country it most likely comes from, without any country names in the prompt.

Together, the three blocks form a progressive difficulty ladder from recall to application to inverse identification.

This design was motivated by a validity problem identified in preliminary testing. Preliminary tests on the original B1 and B2 sets showed that most models achieved over 90% accuracy. However, these results may reflect memorisation of Meyer’s framework during pre-training rather than genuine

cross-cultural understanding. The benchmark therefore risked measuring recall instead of cultural reasoning.

To address this, the dataset was redesigned so that memorised stereotypes would not be enough. Gold answers are derived from Meyer's numerical scales, less familiar country pairs are included, and counter-stereotype items test reliance on regional assumptions. The benchmark contains three blocks: B1 tests relative cultural positions, B2 their application to workplace misunderstandings, and B3 identification of cultural patterns from behaviour.

To illustrate the three formats, consider the following real items from the dataset, all on the *Communicating* dimension:

B1 (Yes/No): *In a project status meeting, a manager from the United States and a manager from France each explain why a deadline slipped. Is the manager from the United States more likely than the manager from France to state the point explicitly and spell out the details, leaving little to be inferred? Answer only Yes or No.*

Gold answer: Yes — the United States scores 5 vs. France at 58 on the Communicating scale (separation = 53 points).

B2 (Applied MCQ): *The United States emails France: "Feedback required on proposal by end of week." France replies with a lengthy message covering context, relationships, and background before briefly addressing the task. The United States concludes the reply missed the point. Which interpretation best explains both the structure of the reply and the United States's misreading of it?*

A. The high-context communicator embedded the key information in surrounding context rather than stating it directly, while the low-context communicator expected explicit delivery. **[Correct]**

B. The reply was assembled by several contributors with competing priorities, so the key task instruction ended up buried beneath sections added late for unrelated reasons.

C. The responder was managing several urgent deadlines that week and addressed

the request only briefly at the very end of the message.

D. The high-context party was still investing in the personal rapport it treats as a precondition for commitment, and used the reply to strengthen that relationship rather than close out the task.

B3 (Identification MCQ): *During a client handover call, a project manager lists every assumption, constraint, and open question in explicit detail before moving to the next topic. When a client asks a follow-up question, the manager restates the answer using different words to ensure nothing is left to interpretation. A professional from which country is most likely to behave in the way described above?*

A. Japan **B.** Denmark **C.** Italy **D.** the United States **[Correct]**

5.2 B1 Factual Yes/No Questions

Each B1 item asks whether a professional from country A is more likely than one from country B to display a behaviour associated with one Meyer dimension.

The correct answer is calculated from the countries' positions on the corresponding 0–100 scale. Only pairs separated by at least 25 points are used, ensuring clear comparisons.

Items are organised into *pair_groups*, where the same scenario is repeated with different country pairs. This measures consistency across familiar and less documented countries.

Additionally, around 35% of items are **counter-stereotype traps**. These items pair two countries from the same broad macro-region (for example, both from Confucian Asia, or both from Latin Europe), but that sit on *opposite* poles of the dimension being tested. The trap targets a common shortcut: a model that has learned to associate a region with a cultural tendency will generalise that tendency to all countries in the region, even when individual countries within it differ sharply. For example, Japan and China both belong to Confucian Asia, yet Japan is strongly linear-time-oriented while China is flexible-time-oriented on the *Scheduling* dimension; France and Sweden are both European, yet France sits at the confrontational pole of *Disagreeing* while Sweden

strongly avoids open conflict. A model that answers correctly on these pairs cannot be relying on regional shortcuts — it must have learned the dimension at the country level. The accuracy gap between standard items and counter-stereotype items is therefore a direct measure of shortcut learning. Variant A uses one template per dimension, while Variant B uses two to introduce wording variation. Each of the two variants (A and B) contains 320 items (40 per dimension); the two variants share the same country pairs by design, so that comparing model performance across them isolates sensitivity to wording from cultural content.

5.3 B2 Applied MCQ Questions

B2 evaluates whether models can apply cultural knowledge to real situations. Each item describes a workplace misunderstanding between professionals from two countries and asks for the most plausible explanation.

The correct option describes the cultural mechanism, while three distractors propose administrative, interpersonal, or technical causes.

The correct answer option never names the countries involved, describing only the cultural mechanism in abstract terms. This prevents the model from exploiting country-name cues to identify the right answer. The `pair_groups` structure is retained to measure consistency across country combinations. Correct-answer positions are randomised to remove the option-A bias found in the original dataset. Each variant contains 160 items (20 per dimension). Because B2 requires applied reasoning rather than a binary factual judgement, each item carries a scoring weight of 2, compared to a weight of 1 for each B1 item. When computing the final benchmark score, B1 and B2 results are combined as a weighted average: each correct B2 response contributes twice as much as a correct B1 response, reflecting the higher cognitive demand of identifying a culturally-grounded explanation over selecting a Yes/No answer. The weight of 2 was chosen as the minimal integer multiplier that preserves this asymmetry without over-penalising factual errors relative to applied ones.

5.4 B3 Identification MCQ questions

B3 reverses the inference direction of B1 and B2. It presents an observable workplace behaviour and asks which country the person is most likely to come from.

Each scenario represents one pole of one Meyer

dimension and contains no explicit cultural markers. The options include the most extreme country on that pole, a less extreme country on the same side, an opposite-pole country, and a neutral country near the midpoint.

B3 contains 80 items, evenly distributed across the eight dimensions and balanced between low- and high-pole behaviours. Answer positions are randomised, and each item has a weight of 2. Together, the three blocks test directional knowledge, applied explanation, and reverse identification.

6 Benchmark and Validation Results

6.1 External-Corpus Mapping

The mapping of external corpora carries several measured limitations. First, the "Leading" dimension acts as an attractor in the final CCD mapping: when the LLM assigns Leading, it is correct only 47% of the time (7 of 15 against ground truth), while recall remains high (7 of 10). Leading is where the model defaults when an item is ambiguous; the bias is bounded and documented, but not eliminated.

Second, inter-annotator agreement on CCD is moderate ($k = 0.56$), and the Deciding-versus-Leading boundary persists as the dominant source of disagreement even under a pre-registered guideline with answer poles visible — 6 of the 22 adjudicated cases. Agreement on EtiCor++ is higher ($k = 0.80$), consistent with its simpler declarative format.

Third, the first annotation round was conducted non-blind (mapper labels visible to one annotator), and an anchoring pattern from the v1 mapper's labels persisted into a subsequent blind round; only an explicit warning in the round-3 guideline eliminated it. We flag this as a methodological caveat and a general warning: automatic mappings should not be validated by annotators who have seen their output, as the contamination is not merely copied labels but a learned criterion that survives their removal.

Fourth, sample sizes bound the precision of our estimates: 90 ground-truth items for CCD-Bench (78% accuracy carries roughly ± 8 -point interval) and 91 for EtiCor++ (92%, roughly ± 6). The gaps over the rule-based baselines (57% and 92% versus 57%) survive these intervals comfortably, but the point estimates should not be over-read.

Fifth, all LLM mapping labels come from a single annotator model




Figure 1: Overview of the Cultural Bias Benchmark system pipeline, illustrating the flow from raw datasets through multi-model LLM inference to automated evaluation and reporting.

(llama-3.3-70b-versatile); no inter-model agreement was measured. Finally, the EtiCor++ rules mapper never received an explicit GLOBE exclusion list (unlike the final CCD prompt), so while its 92% holds on the audited sample, the Unmapped set was not re-checked for over-inclusion of GLOBE content. The (2.5–5)% reformulation-survival figure for EtiCor++ is an estimate from an exploratory pipeline, not a measured result.

6.2 Benchmark Evaluation

The LLM Cultural Map Benchmark is a systematic framework designed to evaluate cross-cultural awareness in language models across eight behavioral dimensions based on Erin Meyer’s Cultural Map and ten GLOBE cultural cluster divisions (House et al., 2004). To overcome data contamination and prevent models from scoring high purely through memorization, the benchmark introduces structured test blocks—including binary Yes/No comparative items (B1) and applied multiple-choice scenarios (B2), complemented by 35% counter-stereotype items to penalize regional shortcut learning. Results from the evaluation reveal significant disparities in performance, with models like gpt-oss-20b achieving strong global accuracy, while other models display prominent regional blind spots and severe vulnerabilities across specific cultural dimensions. The exact mapping

of the country-based cultures analyzed in Meyer’s framework and the ten GLOBE regions used in our evaluations is published in our GitHub repository.

7 Results

The evaluation of Large Language Models reveals significant cross-cultural misalignment, showing how this type of technology can be inherently biased. Overall performance across all three blocks (B1, B2, and B3) indicates substantial differences among the tested models, with GPT-OSS-20b achieving the best results by a wide margin over the rest. The remaining models score between 58% and 63%, as shown in Table 1, suggesting meaningful but incomplete cross-cultural alignment.

Model	Global Accuracy
gpt-oss-20b	74.31%
llama-3.1-8b	62.36%
command-r	59.86%
llama-3.3-70b	58.33%

Table 1: Accuracy of LLMs on All of the Global LLM Culture Map Datasets (B1, B2 and B3)

When examining regional bias across the diverse cultural clusters, performance fluctuations follow generalized patterns for each model, even though their general capabilities differ. All the evaluated models, as shown in Table 2, even though they differ in their average accuracy, tend to fail mainly in tasks related to Nordic and sub-Saharan countries. They generally overperform in Western countries. This difference is mainly produced by the information present on their training datasets, as these countries are negatively affected by the bias that either are not big enough to be represented on the internet or do not have proper access to it.

Assessing the models across Meyer’s eight cultural dimensions further exposes distinct vulnerabilities, as shown in Figure 2. The Deciding and Disagreeing dimensions show the lowest accuracy across all models. This pattern is consistent with known tendencies in LLMs to avoid taking explicit positions or expressing open disagreement, behaviours that are precisely what these two dimensions require the model to reason about. Whether this reflects a training artefact or a genuine gap in cross-cultural reasoning remains an open question, but the benchmark isolates it as a measurable signal.

Region	command-r-08-2024	llama-3.1-8b-instant	llama-3.3-70b-versatile	gpt-oss-20b
Anglo	63.45%	62.76%	66.21%	77.24%
Confucian Asia	65.19%	59.26%	65.19%	80.74%
Eastern Europe	60.0%	61.54%	60.0%	72.31%
Germanic Europe	63.06%	58.6%	63.06%	80.25%
Latin America	67.23%	68.07%	65.55%	75.63%
Latin Europe	56.55%	54.48%	62.07%	77.93%
Middle East	68.29%	63.41%	62.6%	78.05%
Nordic Europe	57.72%	57.05%	58.39%	72.48%
Southern Asia	65.12%	62.79%	69.77%	76.74%
Sub-Saharan Africa	55.47%	59.38%	57.03%	67.97%

Table 2: Accuracy of LLMs on Global LLM Culture Map Dataset by Region




Figure 2: Results of measuring LLM-culture awareness in the business domain on Meyer’s dimensions

8 Conclusion and Discussion

In this work, the LLM Culture Map benchmark was introduced to address and measure some gaps in cross-cultural awareness within Large Language Models in the business domain. Existing benchmarks fail to capture interaction-level cultural nuances, as we mentioned in our findings; only 21% of CCD-Bench maps to any Meyer dimension, and an estimated 2.5–5% of EtiCor++ would survive reformulation into closed questions that measure interaction-level cross-cultural awareness in the business domain according to Erin Meyer’s eight behavioral dimensions. By manually creating comparative questions and applying workplace scenarios based on the information we have from the existing benchmarks and Erin Meyer’s book, we are now able to successfully disentangle genuine cross-cultural interaction reasoning from the memorization of regional standards.

Our evaluations reveal the severe disparities

in the used models’ capabilities. GPT-OSS-20b achieved the highest global accuracy at 74.31%, maintaining a substantial margin over llama-3.1-8b (62.36%), command-r (59.86%), and llama-3.3-70b (58.33%). Despite the difference in accuracy measures of all these models, they all tend to perform better in WEIRD contexts (Western, Educated, Industrialized, Rich, and Democratic country-based cultures), while consistently failing in Nordic and Sub-Saharan contexts. Moreover, we have seen critical vulnerabilities in the Deciding and Disagreeing dimensions of Meyer’s model, which reflect the "Stochastic Parrot" nature of LLMs (Bender et al., 2021) of tending to agree with the users rather than independently navigating complex decisions or disagreements. This bias happens because of the tendency of an LLM towards "Cultural Flattening" (Yu et al., 2025), which is to overgeneralize dominant regional traits to entirely distinct but geographically proximate minority cul-

tures.

9 Limitations

Despite the contributions of this benchmark, several important limitations should be acknowledged. First, establishing a robust dataset requires navigating the inherent difficulty of finding questions that map the eight behavioral dimensions of Erin Meyer within professional business contexts. Many existing cross-cultural resources do not align with it, necessitating extensive manual filtering and mapping work. Second, technical constraints, particularly regarding token availability and context windows of the evaluated language model. Finally, our experimental scope was intentionally restricted to fully open and freely accessible models. Although this design choice ensures complete transparency and reproducibility, it leaves room to explore how commercial frontier models would perform under identical evaluation criteria.

References

- Emily M. Bender, Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. 2021. [On the dangers of stochastic parrots: Can language models be too big?](#) *Association for Computing Machinery - Digital Library*.
- Esin Durmus, Karina Nguyen, Thomas I. Liao, Nicholas Schiefer, Amanda Askell, Anton Bakhtin, Carol Chen, Zac Hatfield-Dodds, Danny Hernandez, Nicholas Joseph, Liane Lovitt, Sam McCandlish, Orowa Sikder, Alex Tamkin, Janel Thamkul, Jared Kaplan, Jack Clark, and Deep Ganguli. 2024. [Towards measuring the representation of subjective global opinions in language models](#). *arXiv preprint*.
- Ashutosh Dwivedi, Siddhant Shivdutt Singh, and Ashutosh Modi. 2025. [Eticor++: Towards understanding etiquettical bias in llms](#). *arXiv preprint*.
- Robert J House, Paul J Hanges, Mansoor Javidan, Peter W Dorfman, and Vipin Gupta. 2004. *Culture, leadership, and organizations: The GLOBE study of 62 societies*. Sage publications.
- Cheng Li, Mengzhou Chen, Jindong Wang, Sunayana Sitaram, and Xing Xie. 2024. [Culturellm: Incorporating cultural differences into large language models](#). *arXiv preprint*.
- Zhaoming Liu. 2024. [Cultural bias in large language models: A comprehensive analysis and mitigation strategies](#). *Research Gate*.
- Erin Meyer. 2014. *The Culture Map: Breaking Through the Invisible Boundaries of Global Business*. Public Affairs.
- John X. Morris, Chawin Sitawarin, Chuan Guo, Narine Kokhlikyan, G. Edward Suh, Alexander M. Rush, Kamalika Chaudhuri, and Saeed Mahloujifar. 2025. [How much do language models memorize?](#) *arXiv preprint*.
- Oriane Peter and Kate Devlin. 2025. [Decentralising llm alignment: A case for context, pluralism, and participation](#). In *AAAI Publications*.
- Sriyash Poddar, Yanming Wan, Hamish Ivison, Abhishek Gupta, and Natasha Jaques. 2024. [Personalizing reinforcement learning from human feedback with variational preference learning](#). *arXiv preprint*.
- Hasibur Rahman and Hanan Salam. 2026. [Ccd-bench: Probing cultural conflict in large language model decision-making](#). *arXiv preprint*.
- Shibani Santurkar, Esin Durmus, Faisal Ladhak, Cinoo Lee, Percy Liang, and Tatsunori Hashimoto. 2023. [Whose opinions do language models reflect?](#) *Stanford University Human-Centered Artificial Intelligence*.
- Maarten Sap, Dallas Card, Saadia Gabriel, Yejin Choi, and Noah A. Smith. 2019. [The risk of racial bias in hate speech detection](#). *ACL Anthology*.
- Zhivar Sourati, Alireza S. Ziabari, and Morteza Dehghani. 2026. [The homogenizing effect of large language models on human expression and thought](#). *A Cell Press Journal - Trends in Cognitive Sciences*.
- Haeun Yu, Seogyong Jeong, Siddhesh Pawar, Jisu Shin, Jiho Jin, Junho Myung, Alice Oh, and Isabelle Augenstein. 2025. [Entangled in representations: Mechanistic investigation of cultural biases in large language models](#). *CopeNLU*.